Spacer composition, not motif grammar, drives regulatory signals in DNA sequence models

Background

The billboard model holds that enhancer activity depends on which transcription factor binding sites are present—their absolute identity and composition—not on their relative positions, orientations, or spacing [1,2]. This contrarian position challenges the prevailing assumption that motif grammar carries regulatory meaning and that DNA foundation models learn positional syntax. If the billboard model is correct, then what these models have actually learned is absolute context: the compositional features of the sequence, rather than positional context encoding relational arrangement between regulatory elements. We test this hypothesis directly using a spacer-factored grammar sensitivity framework (SF-GSI) that isolates positional context (motif arrangement) from absolute context (spacer composition)—resolving a critical confound in standard vocabulary-preserving shuffles, which perturb both simultaneously and therefore cannot distinguish grammar from composition.

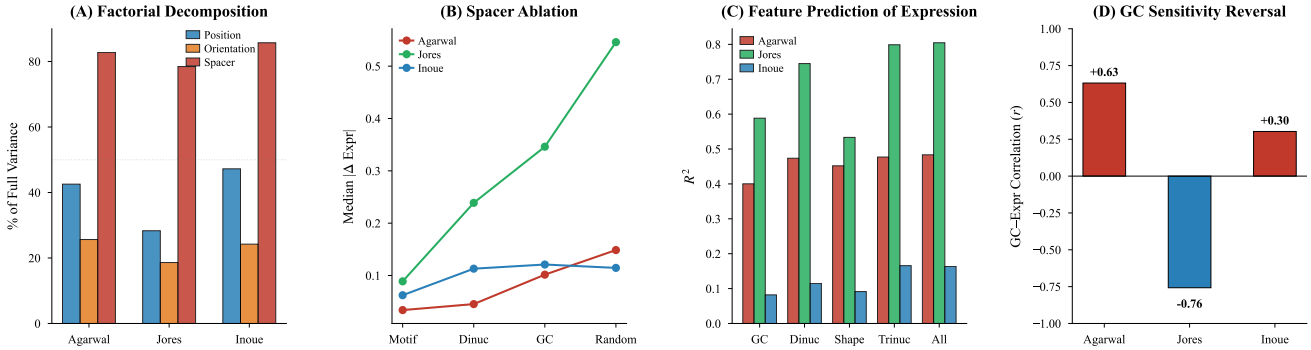

Figure 1: Absolute context dominates apparent grammar effects. (A) Factorial decomposition of prediction variance into position, orientation, and spacer components across three MPRA datasets (DNABERT-2); spacer composition consistently accounts for the largest share, supporting the billboard model. **(B)** Spacer ablation series: replacing spacer DNA (dinucleotide shuffle, GC shift, random replacement) produces larger expression changes than motif-only rearrangement, demonstrating that absolute context—not positional context—drives model sensitivity. **(C)** Feature-based regression (R^2) predicting model expression output from compositional sequence features (GC content, dinucleotide frequencies, DNA shape, trinucleotide counts); compositional features alone explain 40–80% of variance. **(D)** GC–expression correlation reverses sign across datasets, illustrating that models learn dataset-specific compositional biases rather than universal positional grammar.

Methods

SF-GSI computes a grammar sensitivity index, $GSI = \sigma(\hat{y})/|\mu(\hat{y})|$, over vocabulary-preserving motif shuffles and factors the response into positional versus absolute context using matched-spacer perturbations and null-normalized testing. The billboard prediction is concrete: if models encode only absolute context, then controlling for spacer composition should collapse apparent grammar sensitivity to noise—i.e., the residual positional signal should be statistically indistinguishable from a null distribution. We applied SF-GSI to five MPRA datasets (Agarwal, Klein, Inoue, Vaishnav, and Jores), totaling 7,650 model–enhancer measurements: 3 foundation models $\times$ 5 datasets $\times$ 500 enhancers, plus Enformer on 3 human datasets $\times$ 50 enhancers. The model set spans transformer, state-space, and MPRA-native architectures—DNABERT-2, Nucleotide Transformer v2, HyenaDNA, Enformer, and PARM [4–7]—ensuring that conclusions are not artifacts of a single model family. Code: `https://github.com/bryanc5864/GRAMLANG`.

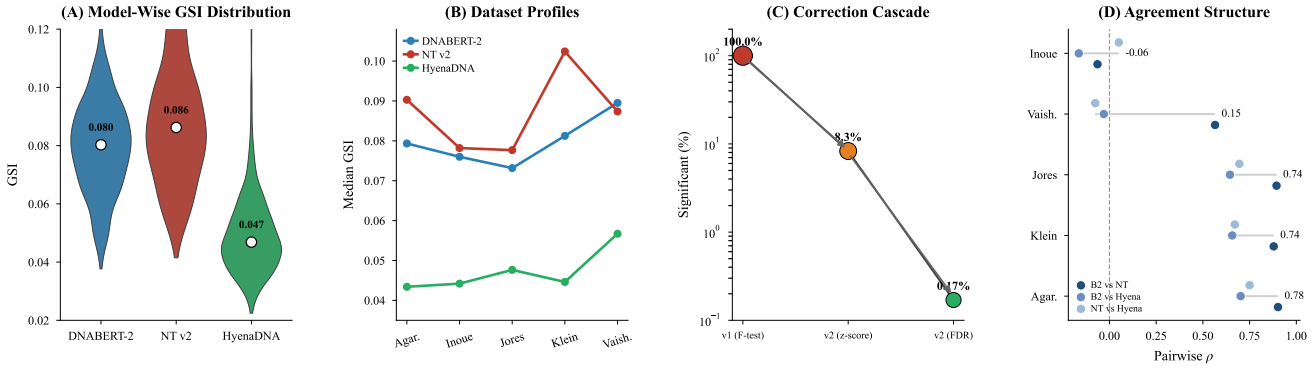

Figure 2: Statistical correction collapses apparent grammar prevalence. (A) Violin plots of raw GSI values by model (DNABERT-2, Nucleotide Transformer v2, HyenaDNA), with median points; all models show similar distributions, indicating that grammar sensitivity estimates are not architecture-specific. **(B)** Median GSI profiles across five datasets for each model; dataset identity drives more variation than model choice, consistent with the billboard view that compositional properties of each library dominate. **(C)** Correction cascade: grammar-sensitive calls fall from 100% under naive thresholding to 8.3% after null normalization and to 0.17% (13/7,650) after FDR correction—a three-order-of-magnitude collapse that matches the billboard prediction. **(D)** Pairwise cross-model GSI correlations by dataset; strong agreement in three datasets but near-zero or negative correlations in Vaishnav and Inoue, showing that residual positional signal is context-dependent rather than reflecting a universal grammar learned across architectures.

Results

Under naive thresholding all tested enhancers appear grammar-sensitive, but this falls to 8.3% after null normalization and to 0.17% (13/7,650) after FDR correction (Fig. 2). That three-order-of-magnitude collapse is the signature predicted by the billboard model: once absolute context is controlled, the apparent prevalence of positional syntax vanishes. Absolute context (spacer composition) accounts for 78–86% of grammar-sensitivity variance (Fig. 1), while positional context—motif-only rearrangements—produces substantially smaller perturbation effects than spacer replacement or GC-shifting perturbations. These models have learned *what* is present, not *where* it is; the signal widely attributed to grammar is, in the vast majority of cases, a compositional signal misread as a positional one.

Positive controls confirm the framework does not erase genuine grammar: with spacer sequence held fixed in synthetic constructs, orientation effects remain strongly detectable across models ($p \leq 7.2 \times 10^{-15}$, with two models reaching numerical underflow). However, these controlled cases are rare exceptions, not the rule. Variance decompositions assign only 0–1.6% of explained variance to arrangement versus 8–22% to motif vocabulary—consistent with the billboard prediction that absolute identity, not positional syntax, is the dominant learned feature. Cross-species transfer is near zero, further suggesting that whatever positional logic models encode is not generalizable. Cross-model agreement concentrates in specific biological contexts rather than reflecting architecture-general positional learning, indicating that dataset provenance matters more than model design.

Conclusions

These results provide direct quantitative support for the billboard model of enhancer regulation as learned by DNA foundation models: these models encode absolute context—what binding sites and sequence features are present—but show negligible encoding of positional context once composition is controlled. The implications extend beyond MPRA grammar analysis: benchmark perturbations that fail to control for absolute context can systematically overstate the positional knowledge that foundation models have acquired. Claims of learned regulatory grammar should be made only under perturbation designs that explicitly factor out absolute context; without that control, compositional signals masquerade as positional ones.

References: [1] Arnosti & Kulkarni, *J. Cell. Biochem.* 94:890, 2005. [2] Long et al., *Cell* 167:1170, 2016. [3] Castro-Mondragon et al., *Nucleic Acids Res.* 52:D142, 2024. [4] Zhou et al., *bioRxiv* 2023, DNABERT-2. [5] Nguyen et al., *bioRxiv* 2023, HyenaDNA. [6] Avsec et al., *Nature Methods* 18:1196, 2021, Enformer. [7] Agarwal et al., *Science* 380:eadh0559, 2023; Jores et al., *Nat. Plants* 7:842, 2021.